

Relational ALGEBRA

↳ basic set of operations for relational model

UNARY OPERATIONS

↳ Select $\sigma_{\text{selection condition}}$

LIST \leftarrow ↳ Project $\pi_{\text{projection list}}$

↳ Rename $\rho_{\text{new name}}$

SET THEORY OPERATIONS

↳ Union \cup

↳ Intersection \cap

↳ Difference $-$

↳ Cartesian Product \times

need to be
TYPE COMPATIBLE

BINARY OPERATIONS

↳ JOIN \bowtie

↳ Division \div $\rightarrow \text{all, every, each}$

ADDITIONAL OPERATIONS

↳ OUTER JOINS

↳ OUTER UNION

↳ Aggregate function



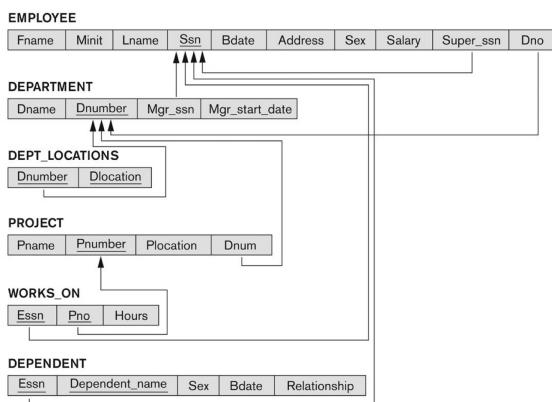
TYPE COMPATIBLE

↳ same no of attributes

↳ same domain for attributes

$$\text{dom}(A_i) = \text{dom}(B_j) \text{ for } i=1,2,\dots,n$$

The set of operations including SELECT σ , PROJECT π , UNION \cup , DIFFERENCE $-$, RENAME ρ , and CARTESIAN PRODUCT \times is called a *complete set* because any other



UNARY OPERATIONS

1. Select (δ) $\delta <\text{condition}_1>, <\text{condition}_2>, \dots, <\text{condition}_n>$ (R)

↳ filter out tuples

- Select the EMPLOYEE tuples whose department number is 4:

$\delta_{DNO=4}$ (EMPLOYEE)

- Select the employee tuples whose salary is greater than \$30,000:

$\delta_{SALARY > 30,000}$ (EMPLOYEE)

2. Project (π) $\pi <\text{attribute list}>$ (R)

↳ keeps certain columns

↳ removes duplicate tuples

P

π

δ

Q) Select fname, lname, salary From Employee E where Dno = 5 AND salary > 300k

1. Single relational algebra expression

$\pi_{fname, lname, salary}(\delta_{DNO=5 \text{ AND } SALARY > 300k}$ (EMPLOYEE))

2. Immediate result relations expression

DEP5_Emp $\leftarrow \delta_{DNO=5 \text{ AND } SALARY > 300k}$ (EMPLOYEE)

final result $\leftarrow \pi_{fname}$ (DEP5_Emp)

3. Rename (P) $P_S(B_1, B_2, \dots, B_n)$ (R)

↳ changes relation name to S
↳ changes column names to B_1, B_2, \dots, B_n

$P(A_1, A_2)$ (EMPLOYEE (A₁, A₂)) \rightarrow A renamed to A₁,
B renamed to A₂

- If we write:

. RESULT $\leftarrow \pi_{FNAME, LNAME, SALARY}$ (DEP5_EMPS)
. RESULT will have the same attribute names as DEP5_EMPS (same attributes as EMPLOYEE)

- If we write:

. RESULT (F, M, L, S, B, A, SX, SAL, SU, DNO) \leftarrow
P RESULT (F, M, L, S, B, A, SX, SAL, SU, DNO) (DEP5_EMPS)

. The 10 attributes of DEP5_EMPS are renamed to F, M, L, S, B, A, SX, SAL, SU, DNO, respectively

Note: the \leftarrow symbol is an assignment operator

A ₁	B ₁	C	D
A ₂	B ₂		

Set Theory Operations

1. Union (U) R ∪ S

- ↳ combine R and S → no of columns must be same
- ↳ no duplicates
- ↳ must be type compatible

- Example:
 - To retrieve the social security numbers of all employees who either work in department 5 (RESULT1 below) or directly supervise an employee who works in department 5 (RESULT2 below)
 - We can use the UNION operation as follows:

$$\begin{aligned} \text{DEP5_EMPS} &\leftarrow \sigma_{\text{DNO}=5}(\text{EMPLOYEE}) \\ \text{RESULT1} &\leftarrow \pi_{\text{SSN}}(\text{DEP5_EMPS}) \\ \text{RESULT2}(\text{SSN}) &\leftarrow \pi_{\text{SUPERSSN}}(\text{DEP5_EMPS}) \\ \text{RESULT} &\leftarrow \text{RESULT1} \cup \text{RESULT2} \end{aligned}$$
 - The union operation produces the tuples that are in either RESULT1 or RESULT2 or both

Some properties of UNION, INTERSECT, and DIFFERENCE

- Notice that both union and intersection are commutative operations; that is
 - $R \cup S = S \cup R$, and $R \cap S = S \cap R$
- Both union and intersection can be treated as n-ary operations applicable to any number of relations as both are associative operations; that is
 - $R \cup (S \cup T) = (R \cup S) \cup T$
 - $(R \cap S) \cap T = R \cap (S \cap T)$
- The minus operation is not commutative; that is, in general
 - $R - S \neq S - R$

2. Intersection (n) R ∩ S

- ↳ common tuples of R and S
- ↳ must be type compatible

3. Difference - R - S

- ↳ all tuples that are in R but not S
- ↳ must be type compatible

4. Cartesian Product × R(A₁, A₂, ..., A_n) × S(B₁, B₂, ..., B_n)

- ↳ combine tuples from 2 diff relations
- ↳ dont need to be type compatible

- Generally, CROSS PRODUCT is not a meaningful operation
 - Can become meaningful when followed by other operations
- Example (not meaningful):
 - $\text{FEMALE_EMPS} \rightarrow \sigma_{\text{SEX}='F'}(\text{EMPLOYEE})$
 - $\text{EMPNAMES} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SSN}}(\text{FEMALE_EMPS})$
 - $\text{EMP_DEPENDENTS} \leftarrow \text{EMPNAMES} \times \text{DEPENDENT}$
- EMP_DEPENDENTS will contain every combination of EMPNAMES and DEPENDENT
 - whether or not they are actually related

- To keep only combinations where the DEPENDENT is related to the EMPLOYEE, we add a SELECT operation as follows
- Example (meaningful):
 - $\text{FEMALE_EMPS} \leftarrow \sigma_{\text{SEX}='F'}(\text{EMPLOYEE})$
 - $\text{EMPNAMES} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SSN}}(\text{FEMALE_EMPS})$
 - $\text{EMP_DEPENDENTS} \leftarrow \text{EMPNAMES} \times \text{DEPENDENT}$
 - $\text{ACTUAL_DEPS} \leftarrow \sigma_{\text{SSN}=\text{ESSN}}(\text{EMP_DEPENDENTS})$
 - $\text{RESULT} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{DEPENDENT_NAME}}(\text{ACTUAL_DEPS})$
- RESULT will now contain the name of female employees and their dependents

÷ = all, every
 U = OR
 does not = —
 ∩ = common = either/or

6:50 AM Tue 19 Dec ...

Division	Course
Sid	Cid
S1	C1
S2	C1
S1	C2
S3	C2

A(X,y) / B(y) = H means \times values for that that should be tuple $\langle x,y \rangle$ for every y value of relation E(Sid, Cid) / C(Cid) = S(?)

student enrolled in every course

Binary Operations



1. Division \div $R(A) \div S(B)$

\hookrightarrow applied on 2 relations

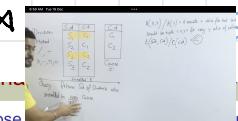
- DIVISION Operation

- The division operation is applied to two relations
- $R(Z) \div S(X)$, where X subset Z . Let $Y = Z - X$ (and hence $Z = X \cup Y$); that is, let Y be the set of attributes of R that are not attributes of S .
- The result of DIVISION is a relation $T(Y)$ that includes a tuple t if tuples t_R appear in R with $t_R[Y] = t$, and with $t_R[X] = t_s$ for every tuple t_s in S .
- For a tuple t to appear in the result T of the DIVISION, the values in t must appear in R in combination with every tuple in S .

$$R \div S \\ R | \leftarrow \pi_{SSN} (\text{EMPLOYEE})$$

2. JOIN \bowtie

Binary Relations



name of the

- Example: Suppose we want to find the manager of each department.

- To get the manager's name, we need to combine each DEPARTMENT tuple with the EMPLOYEE tuple whose SSN value matches the MGRSSN value in the department tuple.
- We do this by using the join \bowtie operation.
- $\text{DEPT_MGR} \leftarrow \text{DEPARTMENT} \bowtie_{\text{MGRSSN}=\text{SSN}} \text{EMPLOYEE}$
- $\text{MGRSSN}=\text{SSN}$ is the join condition
- Combines each department record with the employee who manages the department
- The join condition can also be specified as $\text{DEPARTMENT.MGRSSN} = \text{EMPLOYEE.SSN}$

Equi Join $R \bowtie_{\text{join condition}} S$

Natural Join $R \bowtie_{\text{conditions}} S$

- Example: To apply a natural join on the DNUMBER attributes of DEPARTMENT and DEPT_LOCATIONS, it is sufficient to write:
- $\text{DEPT_LOC} \leftarrow \text{DEPARTMENT} \bowtie \text{DEPT_LOCATIONS}$
- Only attribute with the same name is DNUMBER
- An implicit join condition is created based on this attribute: $\text{DEPARTMENT.DNUMBER} = \text{DEPT_LOCATIONS.DNUMBER}$

- Another example: $Q \leftarrow R(A,B,C,D) * S(C,D,E)$
 - The implicit join condition involves each pair of attributes with the same name, AND'd together:
 - $R.C=S.C$ AND $R.D=S.D$
 - Result keeps only one attribute of each such pair:
 - $Q(A,B,C,D,E)$

- Q1: Retrieve the name and address of all employees who work for the 'Research' department.

$\text{RESEARCH_DEPT} \leftarrow \sigma_{\text{DNAME} = \text{'Research'}} (\text{DEPARTMENT})$
 $\text{RESEARCH_EMPS} \leftarrow (\text{RESEARCH_DEPT} \bowtie_{\text{DNUMBER}=\text{DNUMBER}} \text{EMPLOYEE})$
 $\text{RESULT} \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{ADDRESS}} (\text{RESEARCH_EMPS})$

- Q6: Retrieve the names of employees who have no dependents.

$\text{ALL_EMPS} \leftarrow \pi_{\text{SSN}} (\text{EMPLOYEE})$
 $\text{EMPS_WITH_DEPS(SSN)} \leftarrow \pi_{\text{ESSN}} (\text{DEPENDENT})$
 $\text{EMPS_WITHOUT_DEPS} \leftarrow (\text{ALL_EMPS} - \text{EMPS_WITH_DEPS})$
 $\text{RESULT} \leftarrow \pi_{\text{LNAME}, \text{FNAME}} (\text{EMPS_WITHOUT_DEPS} * \text{EMPLOYEE})$

As a single expression, these queries become:

- Q1: Retrieve the name and address of all employees who work for the 'Research' department.

$\pi_{\text{Fname}, \text{Lname}, \text{Address}} (\sigma_{\text{Dname} = \text{'Research'}} (\text{DEPARTMENT} \bowtie_{\text{Dnumber}=\text{Dno}} \text{EMPLOYEE}))$

- Q6: Retrieve the names of employees who have no dependents.

$\pi_{\text{Lname}, \text{Fname}} ((\pi_{\text{SSN}} (\text{EMPLOYEE}) - \rho_{\text{SSN}} (\pi_{\text{ESSN}} (\text{DEPENDENT}))) * \text{EMPLOYEE})$

- Grouping attribute placed to left of symbol
- Aggregate functions to right of symbol
- $\text{DNO } \text{COUNT } \text{SSN, AVERAGE } \text{Salary } (\text{EMPLOYEE})$

Use of the Aggregate Functional operation \mathcal{F}

- $\mathcal{F}_{\text{MAX Salary}} (\text{EMPLOYEE})$ retrieves the maximum salary value from the EMPLOYEE relation
- $\mathcal{F}_{\text{MIN Salary}} (\text{EMPLOYEE})$ retrieves the minimum salary value from the EMPLOYEE relation
- $\mathcal{F}_{\text{SUM Salary}} (\text{EMPLOYEE})$ retrieves the sum of the salary from the EMPLOYEE relation
- $\mathcal{F}_{\text{COUNT SSN, AVERAGE Salary}} (\text{EMPLOYEE})$ computes the count (number) of employees and their average salary
 - Note: count just counts the number of rows, without removing duplicates

Database Design Process

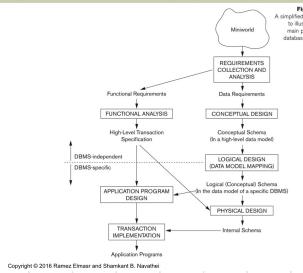
↳ Two main activities

↳ Database Design

↳ Application Design

↳ focuses on programs and interface
that access the DB

Overview of Database Design Process



Copyright © 2010 Ramez Elmasri and Shamkant B. Navathe

Min/Max cardinality

3NF

ERT

NO Relations

Entity Relationship Diagram (ERD)

conceptual view

Entity Types

- ↳ Strong Primary key independent
- ↳ Weak no Primary key dependent on strong entity

e.g. Employee SSN ?
e.g. name birth date ?



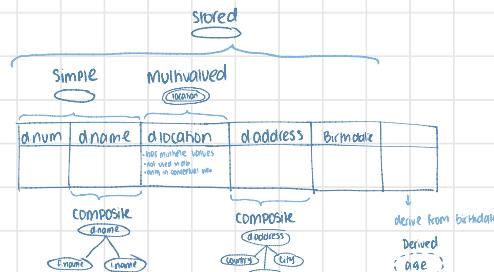
Value Sets

- ↳ specifies set of values

e.g. name has string upto 20 characters
date has value mm-dd-yyyy

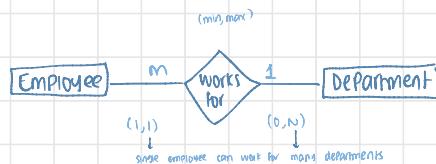
Attributes

- ↳ Simple vs Multivalued
- ↳ Simple vs Composite
- ↳ Stored vs Derived
- ↳ Nested Composite and Multivalued
- ↳ Key attribute



Relationships/Associations

- ↳ 1 : Many (N)
- ↳ Many (N) : 1
- ↳ Many (N) : Many (N)
- ↳ 1 : 1



Structural Constraints

- ↳ Participation
- ↳ Full Participation every employee is in a department
- ↳ Partial Participation not every department has an employee
- ↳ Cardinality (min, max)

USE ANY ONE

Partial: 0
Full: 1

3 ways to represent

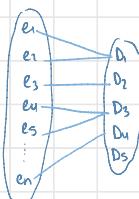
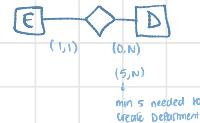
1. Participation



2. Cardinality



3. Structural constraint (min, max)



Corresponding Entity Set

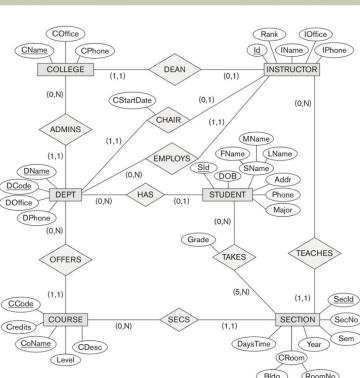


The CAR entity type with two key attributes: Vehicle_id and Year. (a) Entity diagram notation. (b) Entity set with three entities.



© 2016 Ramez Elmasr and Shamkant B. Navathei

Slide 3



© 2016 Ramez Elmasr and Shamkant B. Navathei

Slide 3.



Prelimi
nary
type
of
the
refin

16 Ramez Elmasr and Shamkant B. Navathei



Identifying Relationship

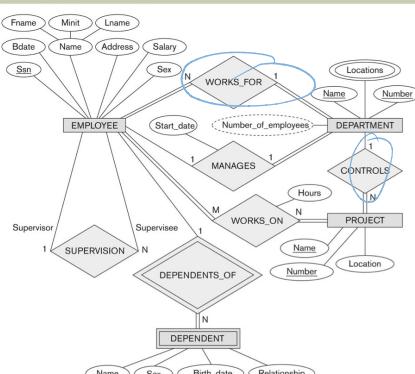


Figure 3.2

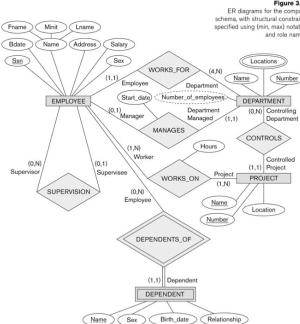
An ER schema diagram for the COMPANY database. The diagrammatic notation is introduced gradually throughout this chapter.

© 2016 Ramez Elmasr and Shamkant B. Navathei

*dependent
↳ primary to decide
NOT PRIMARY KEY*

COMPANY ER Schema Diagram using (min, max) notation

Figure 3.15
ER diagram for the company schema, with structural constraints specified using (min, max) notation and role names.



Copyright © 2016 Ramez Elmasr and Shamkant B. Navathei

Slide 3.

Recursive Relationship Type

↳ same participating entity in different roles

e.g. employee as manager

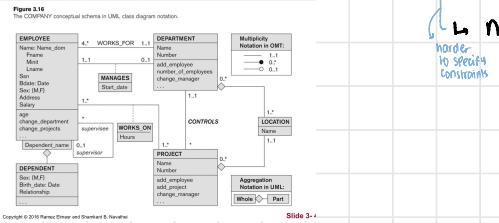
employee as worker

Relationship Type



UML CLASS DIAGRAMS

UML class diagram for COMPANY database schema



Relationship Set

↳ current state of relationship type

Relationships of Higher Degree

↳ Binary : RS types of degree 2

↳ ternary : RS types of degree 3

- L n-ary : RS types of degree $n \geq 2$

Example of a ternary relationship

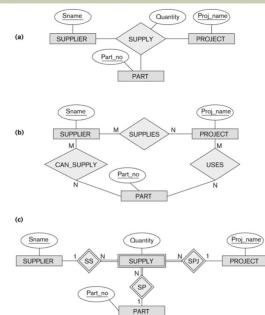


Figure 3.17
Ternary relationship types. (a) The SUPPLY relationship. (b) Three binary relationships not equivalent to SUPPLY. (c) SUPPLY represented as a weak entity type.

Normalization

ch 14
50%

Normalization

- ↳ used on vague data
- ↳ to reduce redundancy
- ↳ null values
 - ↳ attribute not applicable/invalid
 - ↳ attribute value unknown
 - ↳ value known to exist, but unavailable

Relational Database Design

- ↳ grouping of attributes to form good relation schemas
- ↳ it has 2 levels
 1. Logical level → user view
 2. storage level → base relation

ANOMALIES

2. UPDATE ANOMALY

- ↳ chances of redundancy

if EMP.id changes update all
of them in list

3. DELETE ANOMALY

- ↳ chances to have null value

e.g. if we delete a project all emp with
that project will be null
vice versa

1. INSERT ANOMALY

- ↳ chances to have null value
 - e.g. inserting Ali but there's no project yet
- ↳ chances of redundancy
 - e.g. inserting Ali multiple times for diff projects

EXAMPLE OF AN UPDATE ANOMALY

- Consider the relation:
 - EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)
- Update Anomaly:**
 - Changing the name of project number P1 from "Billing" to "Customer-Accounting" may cause this update to be made for all 100 employees working on project P1.

EXAMPLE OF AN INSERT ANOMALY

- Consider the relation:
 - EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)
- Insert Anomaly:**
 - Cannot insert a project unless an employee is assigned to it.
- Conversely**
 - Cannot insert an employee unless an he/she is assigned to a project.

EXAMPLE OF A DELETE ANOMALY

- Consider the relation:
 - EMP_PROJ(Emp#, Proj#, Ename, Pname, No_hours)
- Delete Anomaly:**
 - When a project is deleted, it will result in deleting all the employees who work on that project.
 - Alternately, if an employee is the sole employee on a project, deleting that employee would result in deleting the corresponding project.

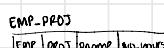
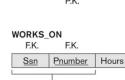
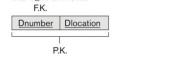
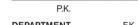
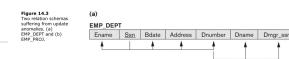


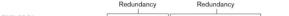
Figure 14.3 Two relation schemas suffering from update anomalies



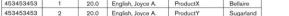
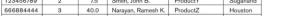
Redundancy



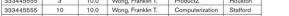
Redundancy Redundancy



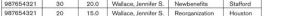
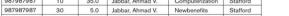
Redundancy Redundancy



Redundancy Redundancy



Redundancy Redundancy



Redundancy Redundancy



Redundancy Redundancy



Redundancy Redundancy

Decompositions

↳ non additive or losslessness of the corresponding join

↳ preservation of the functional dependencies \rightarrow may be sacrificed

↳ very hard
can't be sacrificed

Functional Dependencies

$X \rightarrow Y$ Text \rightarrow course
 X will have diff every text has diff course

	X	Y
	SSN	Emp-name
t_1	123	Ishma

if $t_1(X) = t_2(X)$
then $t_1(Y) = t_2(Y)$

$X = \text{prime}$

$Y = \text{non prime}$
all non prime are determined by prime attributes

X determines Y

e.g.
 $\{SSN \rightarrow ENAME\}$
 $\{P_id \rightarrow \{PNAME, PLOCATION\}\}$
 $\{SSN, PNUMBER \rightarrow HOURS\}$

$\text{Text} \rightarrow \text{course}$ FD

$\text{Teacher} \rightarrow \text{Text}$ NO FD

$\text{Course} \rightarrow \text{Text}$ NO FD

$\text{Teacher} \rightarrow \text{Course}$ NO F

Normalization

↳ lower NF to higher NF \rightarrow normalization form

↳ decomposition

↳ no of decomposition = no of join operations

e.g. 3NF to 4NF

De Normalization

↳ higher NF to lower NF

↳ combine

e.g. 4NF to 3NF

Emp-name	Emp_id	Dep_id	P_no	P_id	P.location
Ali	123	NULL	NULL	NULL	NULL

↓
TO 3NF

Emp-name	Emp_id
Ali	123

Dep_id	Emp_id
NULL	123

P_no	P_id	P.location	Emp_id
NULL	NULL	NULL	123

Superkey

↳ a set of attributes S

↳ such that

$t_1[S] \neq t_2[S]$

Key K

↳ a superkey

↳ such that removal of any attribute from K will cause K to no longer be a superkey

Candidate Key

↳ a relation schema has > 1 keys

1. Primary key

2. Secondary key

↳ Prime attribute: Primary/Secondary key

↳ Non Prime attribute: not a Primary/secondary key

Figure 14.7 Ruling Out FDs

Note that given the state of the TEACH relation, we can say that the FD: Text \rightarrow Course may exist. However, the FDs Teacher \rightarrow Course, Teacher \rightarrow Text and Course \rightarrow Text are ruled out.

TEACH		
Teacher	Course	Text
Smith	Data Structures	Bartram
Smith	Data Management	Martin
Hall	Compilers	Hoffman
Brown	Data Structures	Horowitz

1NF

↳ each cell should have single attribute

↳ no nested relation

↳ no multi valued attributes 

↳ no composite attributes 

↳ can have redundant values → 2NF and > NF
 (partial, irredundant values)

(a) NON 1NF DEPARTMENT			
Dname	Dnumber	Dmgr_ssn	Dlocations
Research	5	333445555	(Bellaire, Sugarland, Houston)
Administration	4	987654321	(Stafford)
Headquarters	1	888665555	(Houston)

(b) SAMPLE DEPARTMENT			
Dname	Dnumber	Dmgr_ssn	Dlocation
Research	5	333445555	Bellaire
Research	5	333445555	Sugarland
Research	5	333445555	Houston
Administration	4	987654321	Stafford
Headquarters	1	888665555	Houston

(c) 1NF with redundancy DEPARTMENT			
Dname	Dnumber	Dmgr_ssn	Dlocation
Research	5	333445555	Bellaire
Research	5	333445555	Sugarland
Research	5	333445555	Houston
Administration	4	987654321	Stafford
Headquarters	1	888665555	Houston

Copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

2NF

↳ no partial dependency

↳ full functional dependency

↳ if every non prime is ↗ on the primary key

AB → C
 composite key
 if C partially depends on AB
 then normalize
 e.g. C depends only on B but not on A

- ↳ $(SSN, PNUMBER) \rightarrow HOURS$ is a full FD since neither SSN → HOURS nor PNUMBER → HOURS hold.
- ↳ $(SSN, PNUMBER) \rightarrow ENAME$ is not a full FD (it is called a partial dependency) since SSN → ENAME also holds.

(a) Normalizing To 2NF EMP PROJ					
Sn	Pnumber	Hours	Ename	Pname	Plocation
FD1					
	FD2				
		FD3			

2NF Normalization					
EPI		EP2		EP3	
Sn	Pnumber	Hours	Ename	Pnumber	Pname
FD1			FD2		FD3

(b) normalizing to 3NF EMP DEPT					
Ename	Sn	Bdate	Address	Dnumber	Dname
ED1				ED2	
					ED3

HOW TO REMOVE MULTIVALUED

1. Decomposition

↳ separate table

↳ make PK FK to the other table

2. Repeat info for id multiple times

EMP_PROJ			Proj
Sn	Ename	Pnumber	Hours
123456789	Smith, John B.	1	32.5
		2	7.5
666884444	Narsen, Remesh K.	3	40.0
454434543	English, Joyce A.	4	20.0
		2	20.0
333445555	Wong, Franklin T.	5	10.0
		3	10.0
		10	10.0
999887777	Zelaya, Alicia J.	20	10.0
		30	30.0
987987987	Jabbar, Ahmad V.	10	35.0
		30	5.0
987654321	Wallace, Jennifer S.	20	20.0
		20	15.0
888665555	Borg, James E.	20	NULL

(b) sample EMP_PROJ		
Sn	Ename	Pnumber
123456789	Smith, John B.	
666884444	Narsen, Remesh K.	
454434543	English, Joyce A.	
333445555	Wong, Franklin T.	
999887777	Zelaya, Alicia J.	
987987987	Jabbar, Ahmad V.	
987654321	Wallace, Jennifer S.	
888665555	Borg, James E.	

Normalizing nested relations at showing nested relations

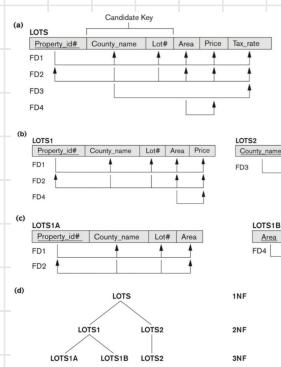
Find PK

↳ if all columns same of any attribute

then that attribute is PK

if can't

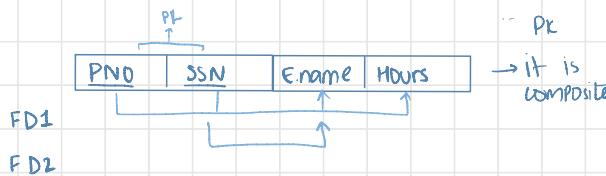
then create a sno column
 give values 1,2,3...n } bad practice
 make it PK



2NF

↳ no Partial dependency

If 2F, no Partial dependency \rightarrow so remove composite



3NF

↳ no transitive dependency



$\text{EMP}(\text{ssn}, \text{emp\#}, \text{salary})$

$\underline{\text{ssn}} \rightarrow \underline{\text{emp\#}} \rightarrow \text{Salary}$
candidate key ✓

$R_1 \rightarrow \underline{\text{ssn}} | \underline{\text{E.name}}$
↓
PK

$R_2 \rightarrow \underline{\text{PNO}} | \underline{\text{ssn}} | \text{hours}$
↓
PK
↓
FK

STEPS

if 1 non prime taken from main table into R1, remove that from main table

R1 and R2 both in 2NF and no partial dependency

every tables non prime is fully dependent on prime

Functional

Transitive Dependency

↳ $X \rightarrow Y \rightarrow Z$

then $X \rightarrow Z$
↳ determines

Functional

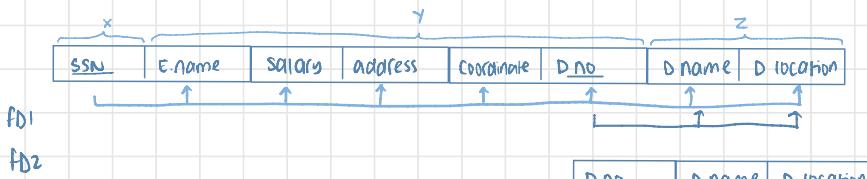
FULL Dependency

Functional

Partial Dependency

↳ Composite PK : 2 or more PK

↳ not 2NF



1NF all attributes depend on a key, Redundancy
2NF all attributes depend on the whole key, No Redundancy, No Partial dependency,
3NF all attributes depend on nothing but the key, No Redundancy, no transitive dependency

1 NF

Employee

EmployeeId	Name	Address	Phone 1
201	Saghir	R288 Karachi	033255
202	Harris	G25 Yorkshire	033543
203	Maxwell	K87 Surrey	035872
204	Andy	Y78 NewCastle	038896
205	Simon	R288 London	038745
206	Sam	F7 Manchester	031210
207	Jim	R88 London	031247
208	Taylor	A4 Manchester	033351

Primary key

Employee		
EmployeeId	Name	Address
201	Saghir	R288 Karachi
202	Harris	G25 Yorkshire
203	Maxwell	K87 Surrey
204	Andy	Y78 NewCastle
205	Simon	R288 London
206	Sam	F7 Manchester
207	Jim	R88 London
208	Taylor	A4 Manchester

Phone		
Phone	Name	
033255	Saghir	
033674	Saghir	
033456	Harris	
035872	Maxwell	
036536	Maxwell	
035972	Andy	
038896	Andy	
038745	Simon	
031210	Sam	
031247	Sam	
033111	Jim	
032266	Jim	
033351	Taylor	
033751	Taylor	

This is 1NF
no repeating columns, no repeating values

?

now this is in
2NF
Both tables have no
partial dependency

2 NF

TEACHER completely depends on COMPOSITE KEY				
StudentId	Course	Name	Marks	Teacher
201	Software-Architecture	Saghir	85	A
202	Software Design	Harris	90	B
201	Quality Assurance	Saghir	75	G
204	English Language	Andy	63	O
205	History	Simon	74	L
206	Project Management	Sam	93	G
205	Software Architecture	Simon	70	A
208	Quality Assurance	Taylor	61	G

COMPOSITE KEY

Result			
StudentId	Course	Name	Marks
201	Software Architecture	Saghir	85
202	Software Design	Harris	90
201	Quality Assurance	Saghir	75
204	English Language	Andy	63
205	History	Simon	74
206	Project Management	Sam	93
205	Software Architecture	Simon	70
208	Quality Assurance	Taylor	61

Teacher	
Course	Teacher
Software Architecture	A
Software Design	B
Quality Assurance	G
English Language	O
History	L
Project Management	G

take away partially dependent to new table

3 NF

MaxMarks transitively depends on Examtype			
StudentId	Name	ExamType	MaxMarks
201	Saghir	Viva	20
202	Harris	Theroy	100
203	Maxwell	Practical	50
204	Andy	Practical	50
205	Simon	Viva	20
206	Sam	Theroy	100
207	Jim	Theroy	100
208	Taylor	Practical	50

Primary key

violating 3NF

Exam		
StudentId	Name	ExamType
201	Saghir	Viva
202	Harris	Theroy
203	Maxwell	Practical
204	Andy	Practical
205	Simon	Viva
206	Sam	Theroy
207	Jim	Theroy
208	Taylor	Practical

Match		
Match#	Teams	Ground
01	Aus v Eng	MCG
02	Zim v Ind	Hobart
03	Sl v Ind	SCG
04	Pak v WI	MCG
05	SA v NZ	Hobart
06	Aus v WI	SCG
07	Pak v Ind	Canberra
08	Eng v Sl	MCG

Marks	
ExamType	MaxMarks
Viva	20
Theroy	100
Practical	50

Capacity	
Ground	Capacity
MCG	80000
Hobart	30000
SCG	65000
Perth	45000
Canberra	100000

v

Single User DBMS

- ↳ one user at a time
- ↳ e.g. home computer

Multi User DBMS

- ↳ many users can access system at a time
- ↳ e.g. airline reservation system, banks, supermarket

Multiprogramming: execute multiple processes concurrently

→ execute some commands of P_1 , then suspend
execute some commands of P_2 , then suspend
resume P_1

Interleaved Processing

- ↳ concurrent execution of processes
- ↳ keep CPU busy, when I/O operation required by switching to another process instead of remaining idle
- ↳ prevents long process from delaying other processes

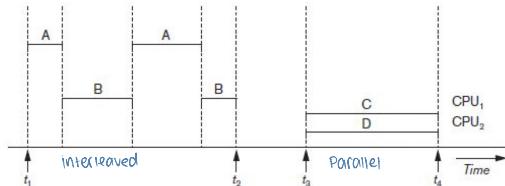


Figure 20.1 Interleaved processing versus parallel processing of concurrent transactions

Transaction

↳ atomic operation
either be completed or not done at all

- ↳ an executing program that forms a logical unit of DB processing
- ↳ has many operations
 1. Insertion
 2. Deletion
 3. Modification
 4. Retrieval

Specifying transaction boundaries

1. begin transaction statement
2. end transaction statement

Types of Transaction

1. Read Only Transaction → only retrieve no update
2. Read Write Transaction → retrieve + update

Parallel Processing

- ↳ uses multiple CPUs
- ↳ parallel processing of multiple processes

Database

- ↳ a collection of named data items

its size is called Granularity
- ↳ data items
 1. DB Record
 2. Disk block
 3. Attribute value of a DB record

larger unit smaller unit

Transaction Processing Systems

- ↳ systems with large databases
- ↳ has many concurrent users
- ↳ requires high availability, fast response time
- ↳ these are independent of item granularity

Read-item(x)

- Reads a database item named X into a program variable named X

Steps:

- Find the address of the disk block that contains item X .
- Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer). The size of the buffer is the same as the disk block size.
- Copy item X from the buffer to the program variable named X .

Write-item(x)

- Writes the value of program variable X into the database item named X
- Process includes finding the address of the disk block, copying to and from a memory buffer, and storing the updated disk block back to disk

- Find the address of the disk block that contains item X .
- Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer).
- Copy item X from the program variable named X into its correct location in the buffer.
- Store the updated disk block from the buffer back to disk (either immediately or at some later point in time).

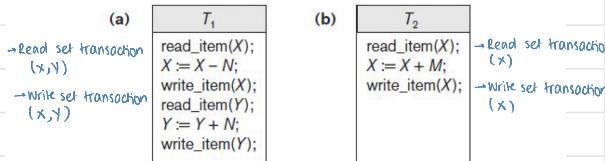


Figure 20.2 Two sample transactions (a) Transaction T_1 (b) Transaction T_2

Concurrency Control

- transactions by multiple users may execute concurrently leading to access and update of same DB items causing an inconsistent DB

Problems

1. LOST UPDATE Problem
2. TEMPORARY UPDATE Problem
3. INCORRECT SUMMARY Problem
4. UNREPEATABLE READ Problem

DBMS Buffers

- DBMS will maintain several main memory data buffers in the database cache
- When buffers are occupied, a buffer replacement policy is used to choose which buffer will be replaced
 - Example policy: least recently used

1. LOST UPDATE PROBLEM → nullifying update of first transaction

- ↳ T_1, T_2 perform interleaved operation on same data item

- ↳ resulting in incorrect value of DB item

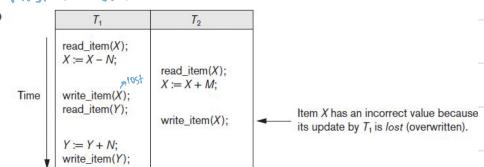


Figure 20.3 Some problems that occur when concurrent execution is uncontrolled (a) The lost update problem

2. DIRTY READ PROBLEM → reading uncommitted data

- ↳ T_1 updates a X and the transaction fails $\xrightarrow{\text{for some reason}}$
 - ↳ Meanwhile updated transaction is read by T_2 before its changed back to its old value
- roll back

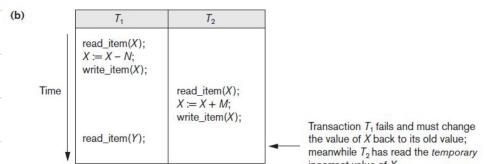


Figure 20.3 (cont'd.) Some problems that occur when concurrent execution is uncontrolled (b) The temporary update problem

3. INCORRECT SUMMARY PROBLEM → invalid result on aggregating data

- ↳ T_1 calculating an aggregate summary function on multiple DB items
- ↳ while other transactions are updating some of these items
- ↳ the aggregate function may calculate some values before and some values after they are updated

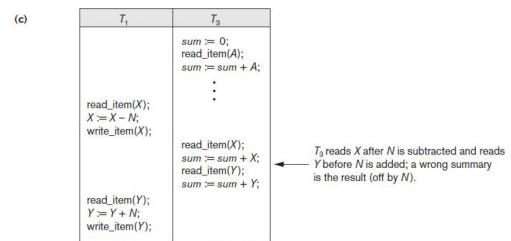


Figure 20.3 (cont'd.) Some problems that occur when concurrent execution is uncontrolled (c) The incorrect summary problem

4. UNREPEATABLE READ PROBLEMS → diff value for same data set variable

- ↳ T_1 reads same DB item twice
- ↳ while T_2 changed its value before T_1 read it the second time
- ↳ so T_1 receive 2 diff values of same item

WHY IS RECOVERY needed

↳ To make sure

1. Committed Transactions: effects are permanently recorded in DB
2. Aborted Transactions: don't affect the DB

Types of Transaction Failures

1. Computer failure, system crash → during transaction execution
hardware/software/network error
2. Transaction/system error → integer overflow, bad casts, logical programming error
3. Local errors → delete item not found
4. Concurrency control enforcement → insufficient account balance in banking
may cause withdraw fund be canceled
↳ aborts transactions
↳ serialization violations
↳ deadlocks
5. Disk failure → block disks lost their data
due to read/write head crash
6. Physical Problems, catastrophes → power conditioning failure
fire, fire, flood, theft

* Failure 1,4 occurs
System must keep sufficient info to recover quickly

common

long recovery time

Transaction States

↳ BEGIN_TRANSACTION

↳ READ/WRITE

↳ END_TRANSACTION

↳ COMMIT_TRANSACTION

↳ ROLLBACK/ABORT

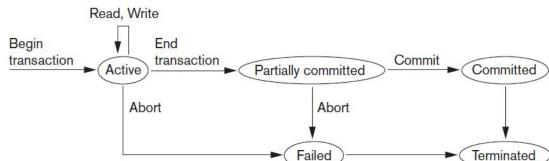


Figure 20.4 State transition diagram illustrating the states for transaction execution

SYSTEM log

↳ to be able to recover from system failures

- ↳ keeps track of transaction operations
- ↳ its a sequential, append only file
- ↳ not affected by failure → except disk/catastrophic
- ↳ Periodically blocked up
- ↳ allows undo/redo operations based on log

log buffer

- ↳ main memory buffers
- ↳ when full, appends to the end of system log file

Commit Point of Transaction

- ↳ when all operations completed successfully
- ↳ and effect of operation stored in log
- Transaction writes a commit record into the log
 - If system failure occurs, can search for transactions with recorded start_transaction but no commit record
- Force-writing the log buffer to disk
 - Writing log buffer to disk before transaction reaches commit point

DBMS Specific Buffer Replacement Policy

- * DBMS Cache: hold disk pages that are currently being processed in main memory buffers
- * Domains: DBMS cache divided into domains

- Page replacement policy
 - Selects particular buffers to be replaced when all are full
- Domain separation (DS) method → LRU page replacement
 - Each domain handles one type of disk pages
 - Index pages
 - Data file pages
 - Log file pages page replacements for each domain handled via LRU
 - Number of available buffers for each domain is predetermined
- Hot set method → page replacement algo
 - Useful in queries that scan a set of pages repeatedly join operation in nested loops
 ↳ disk pages that will be accessed repeatedly
 - Does not replace the set in the buffers until processing is completed
- The DBMIN method → page replacement policy
 - Predetermines the pattern of page references for each algorithm for a particular type of database operation
 - Calculates locality set using query locality set model (QLSM)

ACID PROPERTIES: Transactions several properties

→ enforced by concurrency control, recovery methods

- Atomicity
 - Transaction performed in its entirety or not at all
 - Consistency preservation
 - ↑ A transaction takes database from one consistent state to another, if there's no interference from any other transaction
 - Isolation
 - ↑ A transaction must not interfere with/by other transactions
 - Durability or permanency
 - Changes must persist in the database of commit transaction.
These changes must not be lost
- ...
...
- Levels of isolation
 - Level 0 isolation does not overwrite the dirty reads of higher-level transactions
 - Level 1 isolation has no lost updates
 - Level 2 isolation has no lost updates and no dirty reads
 - Level 3 (true) isolation has repeatable reads
 - In addition to level 2 properties
 - Snapshot isolation → another type of isolation
 - ↑ no phantom records

Schedule / History

- Schedule or history of n transactions
 - Order of operations of transactions
 - Operations from different transactions can be interleaved in the schedule
- Total ordering of operations in a schedule
 - For any two operations in the schedule, one must occur before the other

Conflicting operations in a schedule

- Two conflicting operations in a schedule are in conflict, if they satisfy these conditions
 - Operations belong to different transactions
 - Operations access the same item X
 - At least one of the operations is a write item(X)
- Two operations conflict if changing their order results in a different outcome
- Read-write conflict
- Write-write conflict

CHARACTERIZING SCHEDULES BASED ON RECOVERABILITY

RECOVERABLE SCHEDULES → no abort

- ↳ a committed T should never be rolled back
- ↳ recoverable

NON RECOVERABLE SCHEDULES → abort used

- ↳ a committed T is rolled back
- ↳ should not be permitted by DBMS

T ₁	T ₂
r(x)	
w(x)	r(x)
r(y)	
w(y)	w(x)
commit	commit

T ₁	T ₂
r(x)	
w(x)	r(x)
r(y)	
r(y)	w(x)
abort	commit

T ₁	T ₂
r(x)	
w(x)	r(x)
r(y)	
w(x)	w(x)
abort	commit

so now value
of x is no longer
valid

CASCADING ROLLBACK

- ↳ uncommitted T, may need to be rolled back if dirty read
- ↳ This is time consuming

SOLUTION

CASCADELESS SCHEDULES

- ↳ every T only reads items written by committed T

- ↳ read needs to be committed

STRICK SCHEDULES

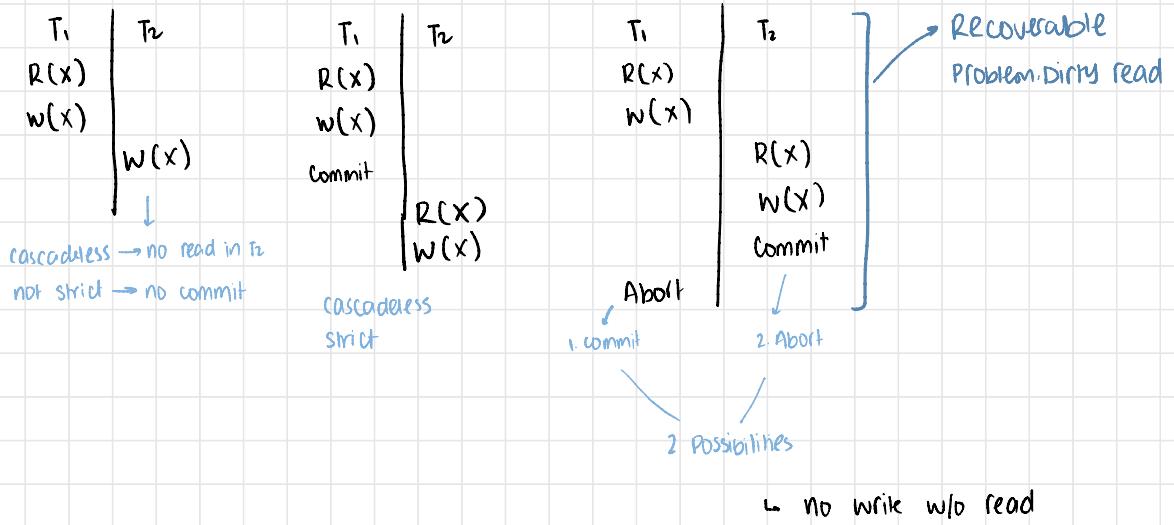
- ↳ every T only reads, writes items written by committed T

- ↳ read, write committed

- ↳ simpler recovery process

→ restore the
before image

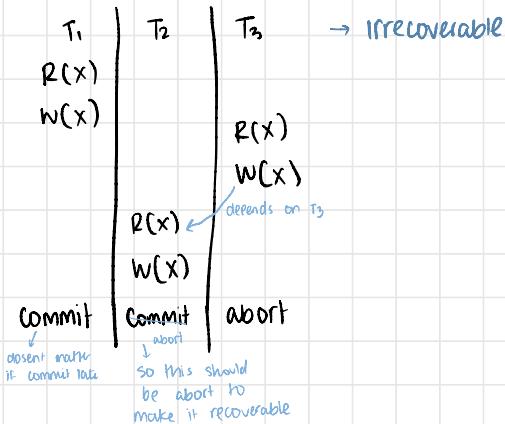
recoverable
cascadeless by default



cascading abort

↳ if T_1 aborts so do the rest of the transaction

↳ concurrency failure transaction



T ₁	T ₂	T ₃
R(x)		
R(z)	R(z)	
		R(x)
		R(y)
w(x)		
Commit		
		w(N)
		Commit
		R(y)
		w(z)
		w(y)
		Commit

cascadeless
strict
recoverable

T ₁	T ₂	T ₃
R(x)		
R(z)	R(z)	
		R(x)
		R(y)
	w(x)	
		w(y)
		Commit
		R(y)
		w(z)
		w(y)
		Commit
		Commit

cascadeless
strict
recoverable

T ₁	T ₂	T ₃
R(x)		
R(z)	R(z)	
		R(x)
		R(y)
w(x)		
Commit		
		w(y)
		Commit
		w(y)
		Commit
		Commit

cascadeless
strict
recoverable

T ₁	T ₂	T ₃
R(x)		
W(x)		
		R(x)
		W(x)
		abort
		R(x)
		W(x)
		Commit
		abort

irrecoverable

abort → now it's recoverable

last transaction
and committed in no
end so it's committed

→ is committed read

→ is committed write

CHARACTERIZING SCHEDULES BASED ON SERIALIZABILITY

Serializable schedules

↳ always correct when concurrent T are executing

↳ places simultaneous T in series $\rightarrow T_1$ before T_2 / T_2 before T_1

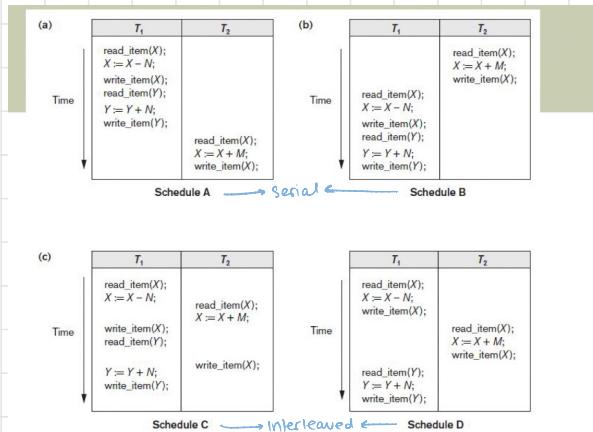


Figure 20.5 Examples of serial and nonserial schedules involving transactions T_1 and T_2 (a) Serial schedule A: T_1 followed by T_2 (b) Serial schedule B: T_2 followed by T_1 (c) Two nonserial schedules C and D with interleaving of operations

Copyright © 2016 Ramez Elmasri and Shamkant B. Navathe

Slide 20-3;

Serial Schedule

↳ T_1 complete, then T_2 start

↳ no concurrency

interleaved schedule

↳ T_1, T_2 concurrency

- Problem with serial schedules
 - Limit concurrency by prohibiting interleaving of operations
 - Unacceptable in practice
 - Solution: determine which schedules are equivalent to a serial schedule and allow those to occur

Serializability

- ↳ interleaved schedule is equivalent to serial schedule
- ↳ interleaved, serial both produce same final state

Conflicting Operations

- ↳ $W(x), R(x)$ no $R(x), R(x)$
- ↳ $W(x), W(x)$
- ↳ $R(x), W(x)$

Conflict equivalence

- ↳ relative order of conflicting operations
is same in both schedules

CHECK Serializability

If conflict equivalence
make precedence graph

↳ make nodes of all transactions

↳ check for conflicting

↳ $W_1(x), R_2(x)$ $T_1 \rightarrow T_2$

↳ $R_1(x), W_2(x)$ $T_1 \rightarrow T_2$

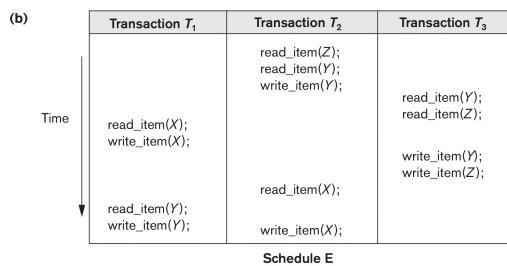
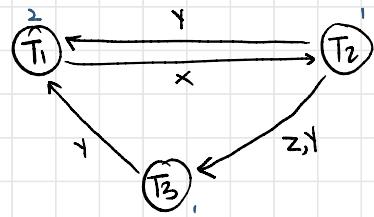
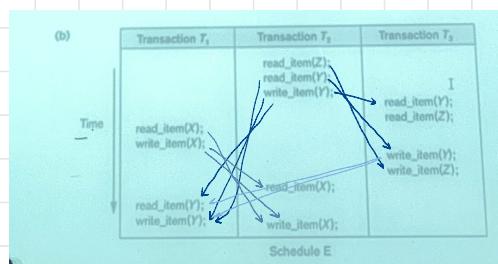
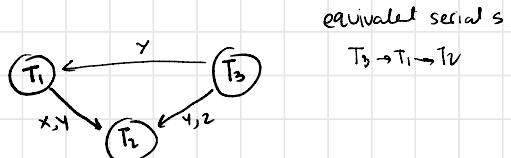
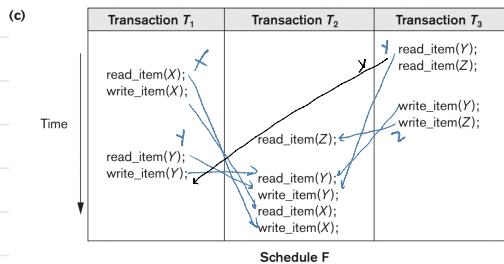
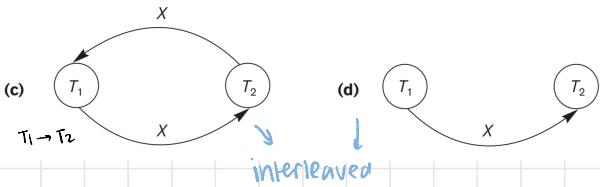
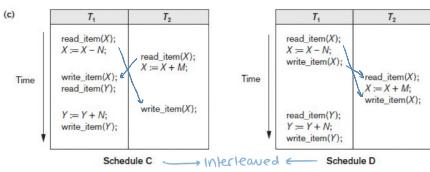
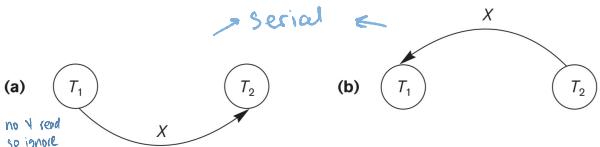
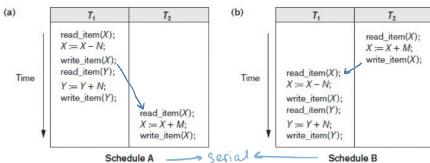
↳ $W_1(x), W_2(x)$ $T_1 \rightarrow T_2$

* min incoming edges runs first

cycle → not equivalent to serial

no cycle → equivalent to serial

Precedence graph



Serializable Schedules

Conflict Serializable

- ↳ has write operation
- ↳ both operations on same data item
- ↳ both operations are by diff T
- Check
 - ↳ draw Precedence graph
 - ↳ if no cycles then ✓
 - ↳ if cycles then X

by default view
resizable

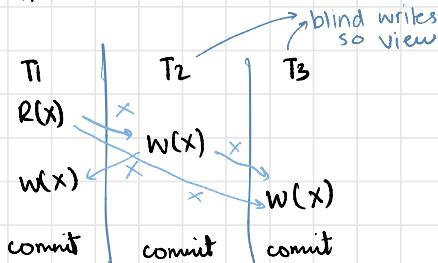
View Serializable

- ↳ if conflict resizable then ✓
- ↳ if
 - ↳ write without a preceding read operation
 - ↳ blind operation ✓
 - ↳ no blind operation X

write

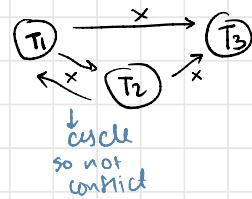
Constrained Write Assumption

- ↳ read operation, then write operation only
- ↳ no blind writes



Unconstrained Write Assumption

- ↳ value written by an operation can be independent of its old value



View Resizable
not conflict Resizable

HOW SERIALIZABILITY IS USED FOR CONCURRENCY CONTROL

SERIALIZABLE SCHEDULE

↳ diff from being serial as

↳ it gives benefit of concurrent execution

without giving up any correctness

- Difficult to test for serializability in practice

- Factors such as system load, time of transaction submission, and process priority affect ordering of operations

- DBMS enforces protocols

- Set of rules to ensure serializability

VIEW EQUIVALENCE OF 2 SCHEDULES

- Same set of T and operations of those T
- read operation value reads result of same write operation in both schedule
write operation produces same results in both schedule
- last operation same in both schedules

OTHER TYPES OF EQUIVALENCE SCHEDULES

Debit Card Transactions

↳ can use non serializable schedules ↳ as order doesn't matter

TRANSACTION SUPPORT IN SQL

- No explicit Begin_Transaction statement
- Every transaction must have an explicit end statement
 - COMMIT
 - ROLLBACK
- Access mode is READ ONLY or READ WRITE
- Diagnostic area size option
 - Integer value indicating number of conditions held simultaneously in the diagnostic area

Isolation level option

- Dirty read
- Nonrepeatable read
- Phantoms

Isolation Level	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No

Table 20.1 Possible violations based on isolation levels as defined in SQL.

Phantoms

- ↳ T₁ reads a set of rows
- ↳ while T₂ inserts a row ↳
- ↳ T₁ is a phantom record as
 - it wasn't there when T₁ started
 - but is there when T₁ ends

SOLUCION

Snapshot Isolation

- ↳ T reads only committed value of data items
- ↳ ensures Phantom record won't occur

CONCURRENCY CONTROL PROTOCOLS

- ↳ set of rules to guarantee serializability

uses

Timestamp

- ↳ unique identifier for each T

TWO PHASE LOCKING PROTOCOL (2PL)

- ↳ locks data items to prevent multiple T from accessing items concurrently
- ↳ has high overhead

MULTIVERSION CONCURRENCY PROTOCOLS

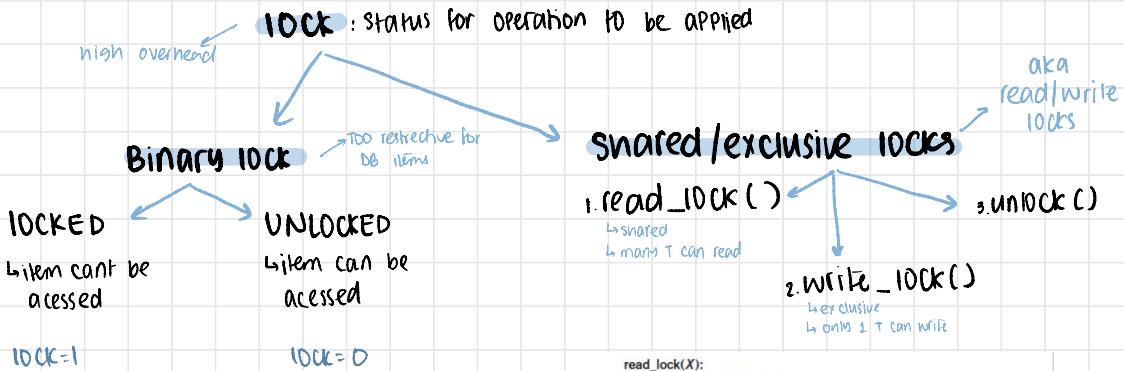
- ↳ uses multiple versions of a data item

LOWER OVERHEAD

OPTIMISTIC PROTOCOLS

- ↳ Validation of a T
- ↳ Certification of a T
- ↳ Then also assume multiple versions of a data item

TWO PHASED LOCKING TECHNIQUES FOR CONCURRENCY CONTROL



- Transaction requests access by issuing a `lock_item(X)` operation

```
lock_item(X):
B: if LOCK(X) = 0           (*item is unlocked*)
    then LOCK(X) ← 1        ("lock the item")
else
begin
    wait (until LOCK(X) = 0
        and the lock manager wakes up the transaction);
    go to B
end;
unlock_item(X):
    LOCK(X) ← 0;           (* unlock the item *)
    if any transactions are waiting
        then wakeup one of the waiting transactions;
```

Figure 21.1 Lock and unlock operations for binary locks

```
read_lock(X):
B: if LOCK(X) = "unlocked"
    then begin LOCK(X) ← "read-locked";
        no_of_reads(X) ← 1
    end
else if LOCK(X) = "read-locked"
    then no_of_reads(X) ← no_of_reads(X) + 1
else begin
    wait (until LOCK(X) = "unlocked"
        and the lock manager wakes up the transaction);
    go to B
end;

write_lock(X):
B: if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
else begin
    wait (until LOCK(X) = "unlocked"
        and the lock manager wakes up the transaction);
    go to B
end;

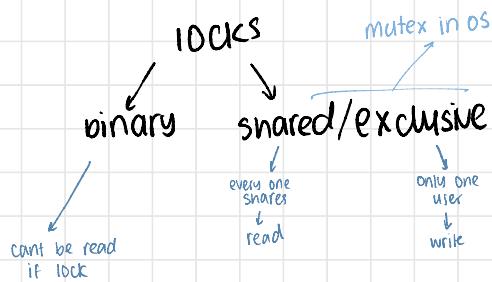
unlock(X):
if LOCK(X) = "write-locked"
    then begin LOCK(X) ← "unlocked";
        wakeup one of the waiting transactions, if any
    end
else if LOCK(X) = "read-locked"
    then begin
        no_of_reads(X) ← no_of_reads(X) - 1;
        if no_of_reads(X) = 0
            then begin LOCK(X) = "unlocked";
                wakeup one of the waiting transactions, if any
            end
    end;
```

LOCK TABLE

- Specifies items that have locks

LOCK manager

- tracks and controls access to locks
- rules enforced by lock manager



read-lock(x):

B:

```
if lock(x) == 'unlock'  
    lock(x) ← 'read-lock'  
    no-of-locks(x) ← 1  
else if lock(x) == "read lock"  
    no-of-reads(x) = no-of-reads(x) + 1  
else begin
```

```
    ask reg-user(x)  
    if ASK == 'NO'  
        wait (until lock(x) == "unlock")  
        "  
    GO TO B
```

END

write-lock(x)

B:

```
if lock(x) == "unlocked"  
    lock(x) ← "write-locked"  
else begin  
    wait (until lock(x) == 'unlocked')  
    GO TO B
```

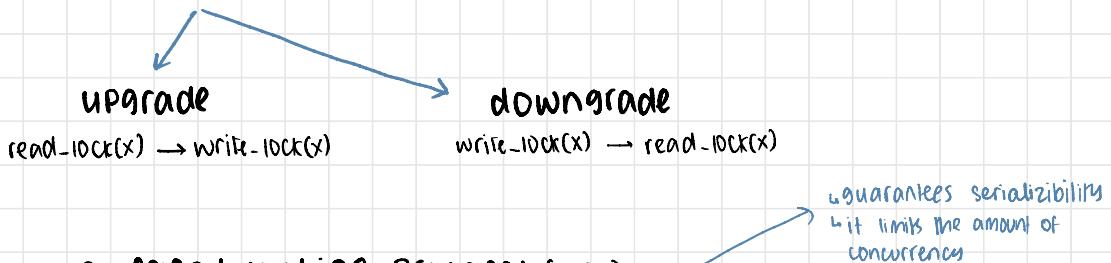
END:

unlock(x):

```
if lock(x) == 'write-lock'  
    lock(x) ← 'unlocked'  
else if lock(x) == "read-lock"  
    no-of-reads(x) = no-of-reads(x) - 1  
    if no-of-reads(x) = 0  
        lock(x) ← 'unlocked'
```

lock conversion

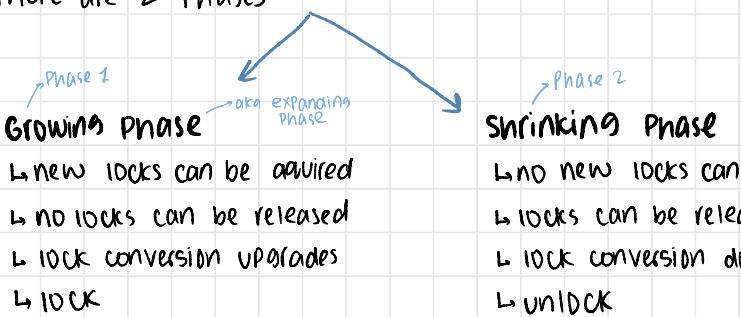
- ↳ a T already holding a lock can convert lock to another state



TWO FACED LOCKING PROTOCOL (TPL)

- ↳ once T releases its first lock, it can no longer acquire new locks

- ↳ There are 2 Phases



* Once gone shrink, can't go back to growing Phase

Figure 21.3 Transactions that do not obey two-phase locking (a) Two transactions T1 and T2 (b) Results of possible serial schedules of T1 and T2 (c) A nonserializable schedule S that uses locks

(a)	T_1	T_2
	<pre>read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); Y := X + Y; write_item(Y); unlock(Y);</pre>

(b)	Initial values: $X=20$, $Y=30$	Result serial schedule T_1 followed by T_2 : $X=50$, $Y=80$	Result of serial schedule T_2 followed by T_1 : $X=70$, $Y=60$
-----	---------------------------------	--	---

(c)	T_1	T_2
Time	<pre>read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre>

	T_1'	T_2'
	<pre>read_lock(Y); read_item(Y); write_lock(X); unlock(Y); read_item(X); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(X); read_item(X); write_lock(Y); unlock(X); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre>

Figure 21.4

Transactions T_1' and T_2' , which are the same as T_1 and T_2 in Figure 21.3 but follow the two-phase locking protocol. Note that they can produce a deadlock.

VARIATIONS OF TWO PHASED LOCKING

BASIC 2PL

↳ as done above

CONSERVATIVE 2PL

↳ T locks all items to be accessed before the T begins

↳ no deadlocks

→ aka static

→ in shrinking phase when it starts

STRICT 2PL

↳ until T commits/aborts

T doesn't release exclusive locks

write

RIGOROUS 2PL

↳ until T commits/aborts

T doesn't release any locks

read, write

→ in expanding phase till it ends

hence no other T can read/write an item that is written by T unless T is committed

CONCURRENCY CONTROL SYSTEM

↳ generates read-lock, write-lock requests

DEADLOCK

↳ set of locked T, each holding a data item and waiting to acquire a data item held by another T

↳ process stuck in circular waiting for the resources

* deadlock implies starvation ↳ starvation does not imply deadlock

STARVATION

↳ when a T waits because it requires a data item to run, that is never allocated to this T

↳

	T_1'	T_2'
Time	<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>write_lock(Z);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>write_lock(Y);</code>

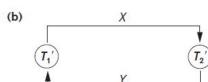


Figure 21.5 Illustrating the deadlock problem (a) A partial schedule of T_1' and T_2' that is in a state of deadlock (b) A wait-for graph for the partial schedule in (a)

Deadlock Prevention Protocols

1

↳ conservative 2PL

↳ lock all data items it needs in advance

2

↳ Order all items in DB

↳ T that needs several items will lock them in that order

both limit concurrency

Wait die

↳ If T_y wants to access item locked by T_o

T_y is aborted and restarted with same priority

↳ If T_o wants to access item locked by T_y

T_o will have to wait

Time Stamp Protocols

$T_y = \text{Younger } T$

$T_o = \text{Older } T$

Wound wait

↳ If T_y wants to access item locked by T_o

T_y is has to wait

↳ If T_o wants to access item locked by T_y

T_o will have access

↳ avoids starvation

No Waiting algo (NW)

↳ If T can't obtain a lock

it is aborted and restarted later
without checking if deadlock
occurs or not

↳ needlessly aborts, restarts
 T

Cautious waiting algo (CW)

↳ If T_1 is not blocked

then T_2 blocked and allowed to wait

↳ If T_1 is blocked

abort T_1

Deadlock Detection

↳ system checks if deadlock state exists

by making a wait-for graph

Victim selection: which T to abort due to deadlock

Timeouts: if a T waits longer than a predefined time, abort the T

Starvation: when T can't proceed for an indefinite time, while other continue normally

↳ solution: first come first serve queue

Concurrency Control Based on Timestamp Ordering (cont'd.)

- **Timestamp**
 - Unique identifier assigned by the DBMS to identify a transaction
 - Assigned in the order submitted
 - Transaction start time
- Concurrency control techniques based on timestamps do not use locks
 - Deadlocks cannot occur

- **Generating timestamps**
 - Counter incremented each time its value is assigned to a transaction
 - Current date/time value of the system clock
 - Ensure no two timestamps are generated during the same tick of the clock
- **General approach**
 - Enforce equivalent serial order on the transactions based on their timestamps

- **Timestamp ordering (TO)**
 - Allows interleaving of transaction operations
 - Must ensure timestamp order is followed for each pair of conflicting operations
- **Each database item assigned two timestamp values**
 - `read_TS(X)`
 - `write_TS(X)`

- **Basic TO algorithm**
 - If conflicting operations detected, later operation rejected by aborting transaction that issued it
 - Schedules produced guaranteed to be conflict serializable
 - Starvation may occur
- **Strict TO algorithm**
 - Ensures schedules are both strict and conflict serializable

- **Thomas's write rule**
 - Modification of basic TO algorithm
 - Does not enforce conflict serializability
 - Rejects fewer write operations by modifying checks for `write_item(X)` operation

?

DATABASE RECOVERY TECHNIQUES

CHP 22

- Recovery process restores database to most recent consistent state before time of failure
- Information kept in system log

TYPICAL RECOVERY STRATEGIES

- Restore backed up copy of DB → when extensive damage
- Identify any changes that may cause inconsistency → when noncatastrophic failure
 - Some operations may require redo

3. DEFERRED UPDATE TECHNIQUE

- When T commits
 - then physically update DB
 - No undo needed
 - Redo might be needed from log
 - No undo/redo also
 - no steal
- Can only be used for short transactions and transactions that change few items
- Buffer space an issue with longer transactions
- Transaction does not reach its commit point until all its REDO-type log entries are recorded in log and log buffer is force-written to disk

4. IMMEDIATE UPDATE TECHNIQUE

- DB might be updated by some operations before T commits
- Operations recorded in logs
- Then operation changes applied on DB
- Undo / redo also
 - Steal / no force strategy

If it crashes after recording some changes
undo operation done in db from log

UNDO AND REDO OPERATIONS

- Idempotent → executing operation multiple times = executing just once
- Entire recovery process should be idempotent

CACHING OF DISK BLOCKS

DBMS cache: a collection of in-memory buffers

cache directory: tracks which DB items are in buffers

flush: cache buffers flushed to make space for new items

If dirty bit = 1 → modified buffer

→ contents written back to disk before flush

If pin-unpin bit = 1 → page can't be written back to disk yet

Main Strategies when pushing a modified buffer back to disk

dirty bit = 1

Inplace updating

- ↳ writes in same db disk location
- ↳ overwrites the changed values

log for recovery
inorder to
undo/redo

Shadowing

- ↳ writes in diff disk location
- ↳ to maintain multiple versions of data items

- Before-image: old value of data item → BFIM
- After-image: new value of data item → AFIM

Write ahead logging

- ↳ ensure before image is recorded in log and
- ↳ log entry is flushed to disk before AFIM overwrites
- UNDOL TYPE log entries → old value → BFIM → before flush image memory
- redolog entries → new value → AFIM → after flush image memory

Steal/no steal, Force/no force

- ↳ rules for when a page from DB cache can be written to disk

NO Steal approach

- ↳ T commits, only then
- ↳ write updated buffer to disk

Steal approach

- ↳ write updated buffer to disk
- ↳ men T commits

NO Force approach

- ↳ T commits, only then
- ↳ write all updated buffer to disk immediately

Force approach

- ↳ write all updated buffer to disk immediately
- ↳ men T commits
- ↳ eliminates need for redo

Typical database systems employ a steal/no-force strategy

- Avoids need for very large buffer space → steal
- Reduces disk I/O operations for heavily updated pages → no force

- Write-ahead logging protocol for recovery algorithm requiring both UNDO and REDO
 - BFIM of an item cannot be overwritten by its after image until all UNDO-type log entries have been force-written to disk
 - Commit operation of a transaction cannot be completed until all REDO-type and UNDO-type log records for that transaction have been force-written to disk

check Point

↳ another type of entry in log

- Taking a checkpoint
 - Suspend execution of all transactions temporarily
 - Force-write all main memory buffers that have been modified to disk
 - Write a checkpoint record to the log, and force-write the log to the disk
 - Resume executing transactions
- DBMS recovery manager decides on checkpoint interval

time needed to
force write may
delay T

solution

Fuzzy checkpointing

- System can resume transaction processing after a begin_checkpoint record is written to the log
- Previous checkpoint record maintained until end_checkpoint record is written

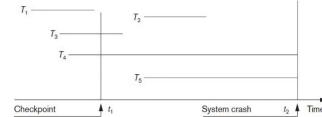


Figure 22.2 An example of a recovery timeline to illustrate the effect of checkpointing

Slide 22-20

Transaction Rollback

- Transaction failure after update but before commit
 - Necessary to roll back the transaction
 - Old data values restored using undo-type log entries
- Cascading rollback
 - If transaction T is rolled back, any transaction S that has read value of item written by T must also be rolled back
 - Almost all recovery mechanisms designed to avoid this

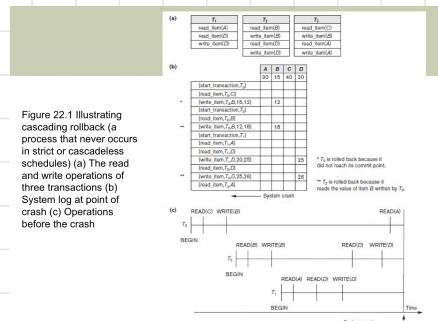


Figure 22.1 illustrating transaction rollback (a) process that never completes in strict or cascaded cursors schedules (a) The read and write operations of three transactions (b) System log at point of crash (c) Operations before the crash

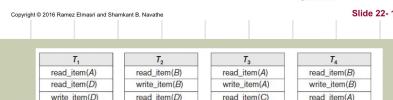


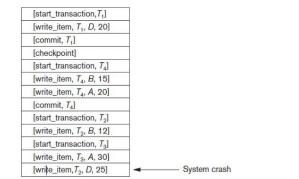
Figure 22.2 An example of a recovery timeline to illustrate the effect of checkpointing

Transactions that Do Not Affect the Database

- Example actions: generating and printing messages and reports
- If transaction fails before completion, may not want user to get these reports
 - Reports should be generated only after transaction reaches commit point
- Commands that generate reports issued as batch jobs executed only after transaction reaches commit point
- Batch jobs canceled if transaction fails

ishma hafeez
notes

refresh
tree



T_1 and T_2 are ignored because they did not reach their commit points.
 T_4 is redone because its commit point is after the last system checkpoint.

Slide 22-1